

GRIP: The Sparks Foundation

Data Science and Business Analytics Intern

Name: Nitha Royals GR

Task 1 :Prediction Using Supervised ML

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
import seaborn as sns
```

```
In [2]: url = "http://bit.ly/w-data"
data= pd.read_csv(url)
```

```
In [3]: print(data.shape)
data.head()
```

Out[3]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [4]: data.describe()
```

Out[4]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [5]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  --
 0   Hours   25 non-null      float64
 1   Scores  25 non-null      int64   
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```

```
In [6]: data.corr(method='pearson')
```

Out[6]:

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

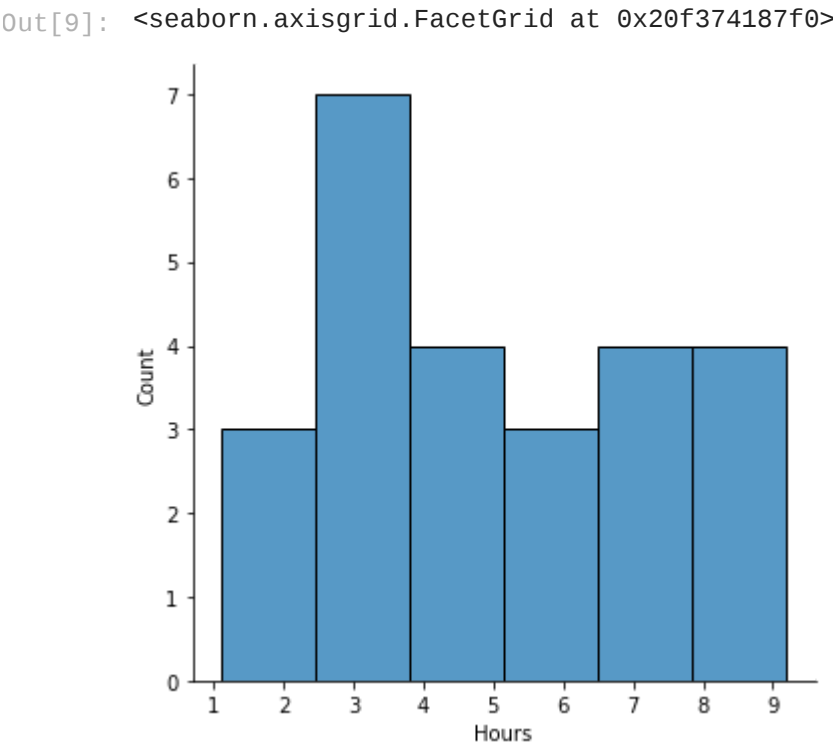
```
In [7]: data.corr(method='spearman')
```

Out[7]:

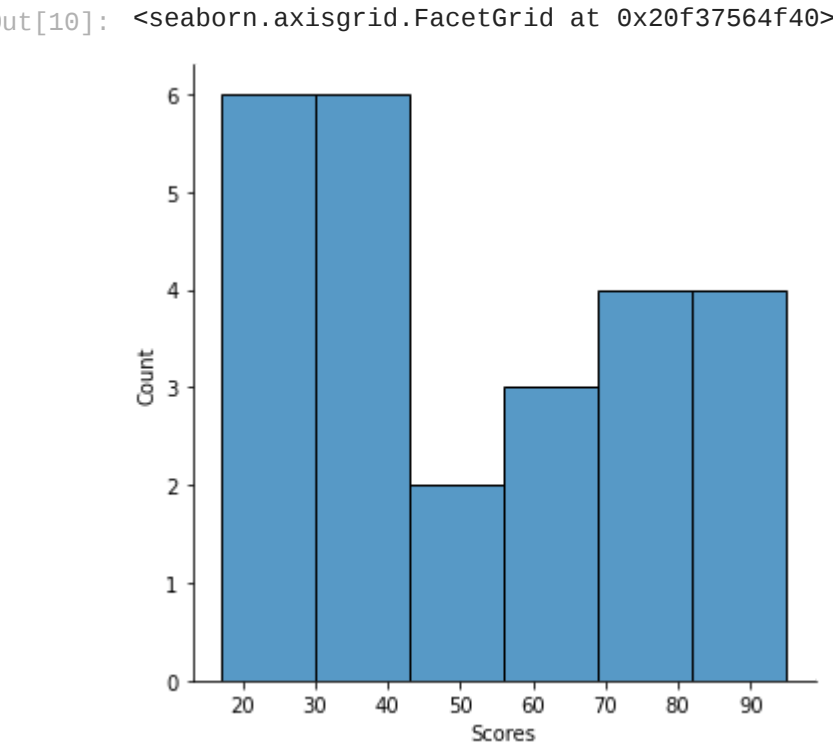
	Hours	Scores
Hours	1.000000	0.971891
Scores	0.971891	1.000000

```
In [8]: hours=data['Hours']
scores=data['Scores']
```

```
In [9]: sns.displot(hours)
```



```
In [10]: sns.displot(scores)
```



Linear Regression


```
In [11]: x = data.iloc[:, :-1].values
y = data.iloc[:, 1].values
```

```
In [12]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=50)
```

```
In [13]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(x_train, y_train)
```

Out[13]: LinearRegression()

```
In [14]: m=reg.coef_
c=reg.intercept_
line=m*x+c
plt.scatter(x, y)
plt.plot(x,line);
plt.show()
```




```
In [15]: y_pred=reg.predict(x_test)
```

```
In [16]: actual_predicted=pd.DataFrame({'Target':y_test,'predicted':y_pred})
actual_predicted
```

Out[16]:

	Target	predicted
0	95	88.211394
1	30	28.718453
2	76	69.020122
3	35	39.273652
4	17	13.365436

```
In [17]: sns.set_style('whitegrid')
sns.displot(np.array(y_test-y_pred))
plt.show()
```



```
In [18]: h=9.25
s=reg.predict([[h]])
print("if a student for {} hpurs per day he/she will score {}% in exam.".format(h,s))
```

if a student for 9.25 hpurs per day he/she will score [91.56986604]% in exam.

```
In [19]: from sklearn import metrics
from sklearn.metrics import r2_score
print('Mean Absolute Error:',metrics.mean_absolute_error(y_test,y_pred))
print('R2 score:',r2_score(y_test,y_pred))
```

Mean Absolute Error: 4.5916495300630285
R2 score: 0.971014141329942